

# SYCL Hands-on Lab – Cambridge Version

Stephen Blair-Chappell  
Intel oneAPI Certified Instructor  
October 2024

# 1. Getting Started

1. Log on to SWIRLES

2. Get list of intel tools via spack or module

```
spack list | grep intel  
module spider oneapi
```

3. Load modules & environment

```
module load intel-oneapi-compilers
```

4. Check versions of compiler

```
icpx --version
```

5. Check GPU Status (only on NVIDIA SYSTEMS)

```
nvidia-smi
```

6. List sycl accelerators available

```
syml-ls
```

# Hands-On Exercise 1 (5 minutes)

1) Use an editor to create code example(ex1.cpp)

2) Compile code

```
icpx -fsycl ex1.cpp -o ex1
```

3) Run program

```
./ex1
```

## QUESTIONS:

- *Which lines of the code runs on the accelerator?*
- *Which Accelerator was used?*

```
#include <iostream>
#include <sycl/sycl.hpp>

int main() {
    try {
        // Step 1: Create a SYCL queue for the default accelerator
        sycl::queue q;

        // Print out the vendor and device name of the selected accelerator
        std::cout << "Running on: "
                  << q.get_device().get_info<sycl::info::device::vendor>()
                  << " "
                  << q.get_device().get_info<sycl::info::device::name>()
                  << std::endl;

        // Step 2: Offload "Hello World" using a lambda function in a kernel
        q.submit([&](sycl::handler &h) {
            // Create a SYCL stream object
            sycl::stream os(1024, 128, h);

            h.single_task([=]() {
                // Print "Hello World" from the accelerator using sycl::stream
                os << "Hello from SYCL!" << sycl::endl;
            });
        }).wait();
    } catch (sycl::exception const &e) {
        std::cout << "An exception occurred: " << e.what() << std::endl;
        return 1;
    }

    return 0;
}
```

# Hands-On Exercise 2(10 minutes)

1) Use an editor to create code example(ex2.cpp)

2) Compile code

```
icpx -fsycl -fsycl-targets=nvptx64-nvidia-cuda  
ex2.cpp -o ex2
```

3) Use the SYCL\_PI\_TRACE to confirm what accelerator the code ran on.

```
SYCL_PI_TRACE=1 ./ex2
```

## QUESTIONS:

- Which line of the code causes code to be associated with gpu?
- Can I change this behaviour at runtime with the ONEAPI\_DEVICE\_SELECTOR ?

```
ONEAPI_DEVICE_SELECTOR="*:cpu" ./ex2
```

- Why did we use the options -fsycl-targets= ?

```
#include <iostream>  
#include <sycl/sycl.hpp>  
  
int main() {  
    try {  
        // Step 1: Create a SYCL queue for gpu accelerator  
        sycl::queue q(sycl::gpu_selector_v);  
  
        // Print out the vendor and device name of the selected accelerator  
        std::cout << "Running on: "  
                    << q.get_device().get_info<sycl::info::device::vendor>()  
                    << " "  
                    << q.get_device().get_info<sycl::info::device::name>()  
                    << std::endl;  
  
        // Step 2: Offload "Hello World" using a lambda function in a kernel  
        q.submit([&](sycl::handler &h) {  
            // Create a SYCL stream object  
            sycl::stream os(1024, 128, h);  
  
            h.single_task([=]() {  
                // Print "Hello World" from the accelerator using sycl::stream  
                os << "Hello from SYCL!" << sycl::endl;  
            });  
        }).wait();  
    } catch (sycl::exception const &e) {  
        std::cout << "An exception occurred: " << e.what() << std::endl;  
        return 1;  
    }  
  
    return 0;  
}
```

# Hands-On Exercise 3(10 minutes)

1) Use an editor to create code example(ex3.cpp)

2) Compile code

```
icpx -fsycl -fsycl-targets=nvptx64-nvidia-cuda  
ex3.cpp -o ex3
```

3) Use the SYCL\_PI\_TRACE to confirm what accelerator the code ran on.

```
SYCL_PI_TRACE=1 ./ex3
```

QUESTIONS:

- Which line of the code causes code to be associated with gpu?
- Can I change this behaviour at runtime with the ONEAPI\_DEVICE\_SELECTOR ?

```
ONEAPI_DEVICE_SELECTOR="*:cpu" ./ex3
```

- What happens if you don't use the options `-fsycl-targets=` ?
- How might I make sure that code running on my laptop gives priority of NVIDIA GPU over Intel GPU

```
[opencl:cpu][opencl:0] Intel(R) OpenCL, 13th Gen Intel(R) Core(TM) i9-13900HX OpenCL 3.0 (Build 0) [2024.18.6.0.02_160000]  
[opencl:gpu][opencl:1] Intel(R) OpenCL Graphics, Intel(R) Graphics [0xa788] OpenCL 3.0 NEO [23.43.27642.52]  
[level_zero:gpu][level_zero:0] Intel(R) Level-Zero, Intel(R) Graphics [0xa788] 1.3 [1.3.27642]~  
[cuda:gpu][cuda:0] NVIDIA CUDA BACKEND, NVIDIA GeForce RTX 4090 Laptop GPU 8.9 [CUDA 12.5]
```

```
#include <sycl/sycl.hpp>  
#include <iostream>  
  
// Define a custom device selection function  
int custom_device_selector(const sycl::device& dev) {  
    // Prioritize GPUs  
    if (dev.is_gpu()) {  
        return 1000; // High score for GPUs  
    }  
    return -1; // Ignore non-GPU devices  
}  
  
int main() {  
    // Create a SYCL queue using the custom device selector function  
    sycl::queue q(custom_device_selector);  
  
    // Print the name of the selected device  
    std::cout << "Running on: "  
                << q.get_device().get_info<sycl::info::device::name>()  
                << std::endl;  
  
    return 0;  
}
```

# Hands-On Exercise 4 & 5 (20 minutes)

1) Use an editor to create code example(**ex4.cpp**) and **fill in the missing code**

2) Compile code

```
icpx -fsycl ex4.cpp -o ex4
```

3) Use the SYCL\_PI\_TRACE to confirm what memory allocations and memory transfers take place

```
SYCL_PI_TRACE=2 ./ex4
```

4) Save a copy of ex4.cpp as **ex5.cpp** and change the code so it uses EXPLICIT memory transfer rather than IMPLICIT memory transfer

5) Compile code

```
icpx -fsycl ex5.cpp
```

6) Use the SYCL\_PI\_TRACE to confirm what memory allocations and memory transfers take place

```
SYCL_PI_TRACE=2 ./ex5
```

## QUESTIONS:

- Explain the difference between the two code examples?
- Is the trace different between the two examples – can you explain why?

```
#include <iostream>
#include <sycl/sycl.hpp>

int main() {
    // Step 1: Create a SYCL queue
    // TODO ;

    // Step 2: Allocate shared memory using USM
    const int N = 16;
    int *data = ;

    // Step 3: Initialize the array on the host
    for (int i = 0; i < N; i++) {
        data[i] = i;
    }

    // Step 4: Submit a kernel to the device to add 5 to each element
    q.parallel_for(N, [=](sycl::id<1> i) {
        /* Add 5 to each element of 'data' */
        //TODO ;
    }).wait();

    // Step 5: Print the updated values on the host
    for (int i = 0; i < N; i++) {
        std::cout << "data[" << i << "] = " << data[i] << std::endl;
    }

    // Step 6: Free the allocated memory
    // TODO ;

    return 0;
}
```

# Hands-On Exercise 6(10 minutes)

1) Use an editor to create code example(ex6.cpp)

2) Compile code

```
icpx -fsycl ex6.cpp -o ex6
```

3) Use the SYCL\_PI\_TRACE to confirm what accelerator the code ran on.

```
SYCL_PI_TRACE=1 ./ex6
```

4) Create dependency graph and convert to png

```
SYCL_PRINT_EXECUTION_GRAPH=always ./ex6
```

```
dot -Tpng <filename> -o <outname>
```

This only works if graphviz is installed

5) Edit the code so the dependencies are implemented another way.

```
#include <sycl/sycl.hpp>
using namespace sycl;
constexpr int N = 42;

int main(int argc, char const *argv[])
{
    queue q;

    // Task A
    auto eA = q.submit([&](handler &h)
    {
        sycl::stream out(1024, 256, h);
        h.parallel_for(N, [=](id<1>)
        { out << "A"; }); });
    eA.wait();

    // Task B
    auto eB = q.submit([&](handler &h)
    {
        sycl::stream out(1024, 256, h);
        h.parallel_for(N, [=](id<1>)
        { out << "B"; }); });

    // Task C
    auto eC = q.submit([&](handler &h)
    {
        sycl::stream out(1024, 256, h);
        h.depends_on(eB);
        h.parallel_for(N, [=](id<1>)
        { out << "C"; }); });

    // Task D
    auto eD = q.submit([&](handler &h)
    {
        sycl::stream out(1024, 256, h);
        h.depends_on({eB,eC});
        h.parallel_for(N, [=](id<1>)
        { out << "D"; }); });

    return 0;
}
```

# Hands-On Exercise 7 (30 minutes)

```
//=====
// Copyright © Intel Corporation
//
// SPDX-License-Identifier: MIT
//=====
#include <sycl/sycl.hpp>
using namespace sycl;

static const int N = 1024;
int main() {
    queue q;
    std::cout << "Device : " << q.get_device().get_info<info::device::name>() << "\n";

    //intialize 2 arrays on host
    int *data1 = static_cast<int *>(malloc(N * sizeof(int)));
    int *data2 = static_cast<int *>(malloc(N * sizeof(int)));
    for (int i = 0; i < N; i++) {
        data1[i] = 25;
        data2[i] = 49;
    }

    // # STEP 1 : Create USM device allocation for data1 and data2
    // # YOUR CODE GOES HERE

    // # STEP 2 : Copy data1 and data2 to USM device allocation
    // # YOUR CODE GOES HERE

    // # STEP 3 : Write kernel code to update data1 on device with square of its value
    q.parallel_for(N, [=](auto i) {

        // # YOUR CODE GOES HERE
    });
}
```

```
// # STEP 3 : Write kernel code to update data2 on device with square of its value
q.parallel_for(N, [=](auto i) {

    // # YOUR CODE GOES HERE
});

// # STEP 5 : Write kernel code to add data2 on device to data1

q.parallel_for(N, [=](auto i) {

    // # YOUR CODE GOES HERE
});

// # STEP 6 : Copy data1 on device to host
// # YOUR CODE GOES HERE

// # verify results
int fail = 0;
for (int i = 0; i < N; i++) if(data1[i] != 12) {fail = 1; break;}
if(fail == 1) std::cout << " FAIL"; else std::cout << " PASS";
std::cout << "\n";

// # STEP 7 : Free USM device allocations
// # YOUR CODE GOES HERE

// # STEP 8 : Add event based kernel dependency for the Steps 2 - 6

return 0;
}
```



# Other SYCL - LABs

- Request Jupyter session
- Clone the oneAPI-Samples
- (everything):

```
git clone https://github.com/oneapi-src/oneAPI-samples.git
```

OR

- Or use oneapi-cl and create new C++ Project based on:

```
oneAPI-samples/DirectProgramming/C++SYCL/Jupyter/oneapi-essentials-training
```

# SYCLOMATIC -LAB

- Clone nvidia examples:

```
git clone https://github.com/NVIDIA/cuda-samples.git
```

- Choose one of the samples to migrate.  
(Not all samples will migrate successfully!)

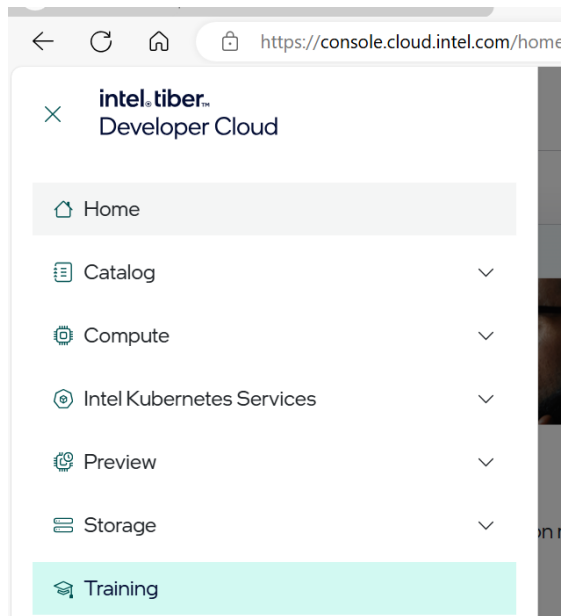
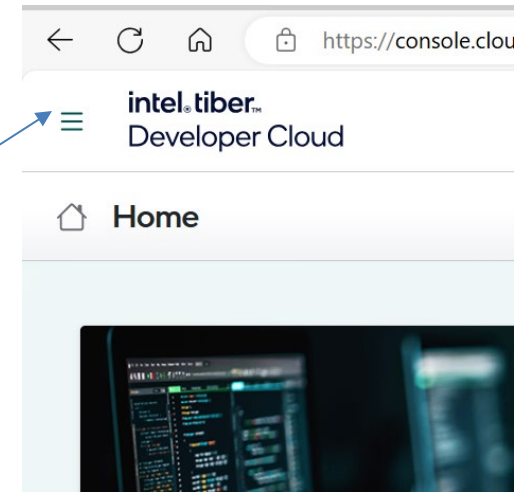
The Intel logo is centered on a blue background. It features the word "intel" in a white, lowercase, sans-serif font. A small, solid blue square is positioned above the first vertical stroke of the letter "i". To the right of the word "intel" is a registered trademark symbol, which consists of the letter "R" enclosed within a circle.

intel®

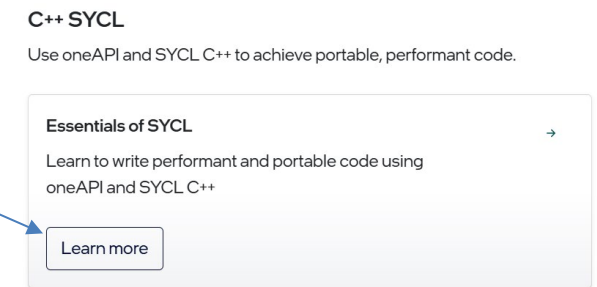
# SYCL Programming Hands-on

Log in to <https://console.cloud.intel.com/>

1. Click here to get the dropdown menu:
2. Choose training:



3. Scroll down and click 'Learn more in 'Essentials of SYCL C++'



4. Click 'Launch Jupyter Notebook'

